

Cursor for Product Managers

An AI-powered platform that automates product discovery for Product Managers by transforming raw customer interviews into actionable product artifacts.

Overview

Cursor for Product Managers helps PMs reduce time spent on manual discovery work. The platform transcribes customer interviews, extracts insights, identifies patterns across conversations, and generates Product Requirements Documents (PRDs) and Jira tickets.

Core Features

- User Authentication – Secure registration and login using JWT
- Project Management – Create and manage multiple product projects
- File Upload – Drag-and-drop support for interview audio/video files
- AI Processing – Multi-agent AI system powered by Gemini 2.0
- Real-time Progress – Live processing updates via Socket.io
- Insights Dashboard – Pain points, quotes, sentiment, and patterns
- PRD Generation – Automatically generated Product Requirements Documents
- Export – Download PRDs or create Jira tickets

Tech Stack

- Frontend: React, Redux Toolkit, Tailwind CSS, React Dropzone, Socket.io-client
- Backend: Node.js, Express, MongoDB, Mongoose, GridFS, JWT
- AI: Gemini 2.0 API with a multi-agent architecture
- Queue: Bull with Redis for background processing
- Real-time: Socket.io
- File Storage: MongoDB GridFS

AI Agents

- Transcriber Agent – Converts audio/video interviews into text with speaker identification
- Insight Agent – Extracts pain points, feature requests, quotes, and sentiment
- Pattern Agent – Identifies trends and common themes across multiple interviews
- PRD Agent – Converts insights into a structured Product Requirements Document

Database Models

User

- email (String, unique)
- password (String, hashed)
- name (String)

- createdAt (Date)

Project

- name (String)
- description (String)
- userId (ObjectId, reference to User)
- createdAt (Date)

Interview

- projectId (ObjectId, reference to Project)
- fileName (String)
- fileId (String, GridFS reference)
- status (uploaded, processing, completed, failed)
- transcript (String)
- insights: painPoints, featureRequests, quotes, sentiment
- createdAt (Date)

PRD

- projectId (ObjectId)
- title (String)
- content: executiveSummary, problemStatement, personas, successMetrics, features
- sourceInterviews (Array of ObjectIds)
- createdAt (Date)

API Endpoints

Authentication

- POST /api/auth/register
- POST /api/auth/login
- GET /api/auth/me

Projects

- GET /api/projects
- POST /api/projects
- GET /api/projects/:id
- PUT /api/projects/:id
- DELETE /api/projects/:id

Interviews

- POST /api/projects/:projectId/interviews/upload
- GET /api/projects/:projectId/interviews
- GET /api/interviews/:id
- DELETE /api/interviews/:id

- GET /api/interviews/:id/insights

PRD

- POST /api/projects/:projectId/prd/generate
- GET /api/prd/:id
- PUT /api/prd/:id

Data Flow

- User uploads interview files which are stored in GridFS
- A Bull job is created and added to the Redis queue
- Worker processes the job and runs the Transcriber Agent
- Transcript is analyzed by the Insight Agent
- Pattern Agent aggregates insights across interviews
- PRD Agent generates a Product Requirements Document on request
- All processing updates are streamed to the frontend via Socket.io